

ShadowCode: Toward (Automatic) External Prompt Injection Attack Against Code LLMs

Yuchen Yang, Yiming Li , *Member, IEEE*, Hongwei Yao , Bingrun Yang , Yiling He, Tianwei Zhang , *Member, IEEE*, Dacheng Tao , *Fellow, IEEE*, and Zhan Qin 

Abstract—Recent advancements have led to the widespread adoption of code-oriented large language models (Code LLMs) for programming tasks. Despite their success in deployment, their security research is left far behind. This paper introduces a new attack paradigm: (automatic) external prompt injection against Code LLMs, where attackers generate concise, non-functional *induced perturbations* and inject them within a victim’s code context. These induced perturbations can be disseminated through commonly used dependencies (e.g., packages or RAG’s knowledge base), manipulating Code LLMs to achieve malicious objectives during the code completion process. Compared to existing attacks, this method does not necessitate control over the model’s training process, unlike backdoor attacks, and can achieve specific malicious objectives that are challenging for adversarial attacks. Furthermore, we propose ShadowCode, a simple yet effective method that automatically generates induced perturbations based on code simulation to achieve effective and stealthy external prompt injection. ShadowCode designs its perturbation optimization objectives by simulating the victim’s code contexts and employs a greedy optimization approach with two enhancement modules: forward reasoning enhancement and keyword-based perturbation design. We evaluate our method across 13 distinct malicious objectives, generating 31 threat cases spanning three popular programming languages. Our results demonstrate that ShadowCode successfully attacks three representative open-source Code LLMs (achieving up to a 97.9% attack success rate) and two mainstream commercial Code LLM-integrated applications (with over 90% attack success

rate) across all threat cases, using only a 12-token non-functional induced perturbation.

Index Terms—Code-oriented large language models, prompt injection attack, code security, AI security.

I. INTRODUCTION

RECENT advancements in generative artificial intelligence (GenAI) have led to its widespread and successful application in code engineering, such as automated programming and code generation [1], [2]. Notably, code-oriented large language models (Code LLMs) such as Codex [3], CodeGemma [4], and CodeGeeX [5] have shown significant promise in code completion.¹ These models efficiently generate complete, functional code from fragmented snippets or natural language prompts, demonstrating superior proficiency in following standardized syntax. This compelling capability has spurred the development of numerous Code LLM-integrated applications, such as GitHub Copilot [1], developed jointly by Github and OpenAI, and has demonstrated notable potential for supporting LLM agents [8].

Despite their impressive performance, recent studies have shown that Code LLMs are susceptible to sophisticated attacks that maliciously manipulate their functional execution. These attacks generally fall into two primary categories: (1) *Backdoor attacks* [9], [10], [11], aim to compromise models by poisoning training data [9], [11] or modifying model parameters [10]. This manipulation causes infected models to generate targeted code snippets when specific, attacker-defined trigger patterns are present in the input. The consequences of such attacks can be diverse, including incorrect code additions, deletions, or modifications [10], as well as the implementation of less secure encryption rules or protocols [9]. (2) *Adversarial attacks* [12], [13] focus on degrading the quality of the generated code by crafting adversarial examples. These examples are typically based on approximated gradients with respect to the model’s inputs and outputs. For instance, Jha et al. [13] demonstrated performance degradation in Code LLMs by altering the function and variable names.

However, we contend that existing attacks possess inherent limitations, hindering their practical threat. Specifically, while backdoor attacks can achieve precise malicious objectives, they necessitate intervention in the model training process, such as modifying training samples [11], [14] or even altering model

Received 25 June 2025; revised 23 May 2026; accepted 2 June 2026. Date of publication 15 June 2026; date of current version 1 September 2026. This work was supported in part by the “Pioneer” and “Leading Goose” R&D Program of Zhejiang under Grant 2026C01020, in part by Kunpeng–Ascend Science and Education Innovation Excellence/Incubation Center, in part by the National Natural Science Foundation of China under Grant 625B1031 and Grant 62441238, in part by the National Natural Science Foundation of China under Grant U2441240, in part by “Ye Qisun” Science Foundation, and in part by the Zhejiang Provincial Natural Science Foundation of China under Grant LD24F020010. (Corresponding author: Yiming Li.)

Yuchen Yang, Hongwei Yao, Bingrun Yang, Yiling He, and Zhan Qin are with the State Key Laboratory of Blockchain and Data Security, Zhejiang University, Hangzhou 310007, China, and also with the Hangzhou High-Tech Zone (Binjiang) Institute of Blockchain and Data Security, Hangzhou 310053, China (e-mail: ychyang@zju.edu.cn; yhongwei@zju.edu.cn; yangbingrun@zju.edu.cn; qinzhan@zju.edu.cn; heyilinge0@gmail.com).

Yiming Li is with the State Key Laboratory of Blockchain and Data Security, Zhejiang University, Hangzhou 310007, China, and also with the College of Computing and Data Science, Nanyang Technological University, Singapore 639798 (e-mail: liyiming.tech@gmail.com).

Tianwei Zhang and Dacheng Tao are with the College of Computing and Data Science, Nanyang Technological University, Singapore 639798 (e-mail: tianwei.zhang@ntu.edu.sg; dacheng.tao@ntu.edu.sg).

The code is available at <https://github.com/LianPing-cyber/ShadowCodeEPI>. This article has supplementary downloadable material available at <https://doi.org/10.1109/TDSC.2026.3703498>, provided by the authors.

Digital Object Identifier 10.1109/TDSC.2026.3703498

¹Code completion currently remains a primary research focus [6], [7] and a foundational paradigm in commercial practice [1], [2], despite the existence of other programming tasks (e.g., code summarization).




Fig. 1. The attack paradigm of external prompt injection involves injecting a concise, non-functional, attacker-generated perturbation into a victim's code context. When the victim user completes the code via a Code LLM, the model is induced by this incorporated attacker-specified code to generate an unsafe output at a specified position, e.g., in `edit_part()` function in this example.

parameters [10]. In real-world scenarios, developers of Code LLM or integrated applications usually acquire the models or datasets from trusted entities (e.g., renowned research institutes or companies), thereby mitigating backdoor threats from their origin. Conversely, existing adversarial attacks [12], [13], despite not requiring training intervention, struggle to achieve precise malicious objectives. This is primarily attributable to their core aim: identifying adversarial samples that degrade code generation quality, which inherently compromises the Code LLM's understanding of the input. These findings raise an intriguing and critical question: *Has the application of Code LLMs already achieved a sufficient level of security in practice?*

Unfortunately, the answer to the above question is negative. Drawing inspiration from prompt injection (PI) attacks (e.g., jailbreak attacks) [15], [16], which leverage carefully designed prompts to force LLMs into generating malicious output, we observe that prompt perturbations introduce a further risk of malicious code completion. In this paper, we introduce a new attack paradigm: external prompt injection against Code LLMs.² As illustrated in Fig. 1, an attacker crafts a concise non-functional *induced perturbation* (e.g., comments) containing malicious instructions (e.g., generating code snippets for data deletion [19]). This perturbation is then inserted into the victim's code context. When the victim requests the Code LLM to perform code completion based on this infected code context, the implanted perturbation will induce the model to respond to the *particular* malicious instructions at *specific* locations. Arguably, attackers can now easily and stealthily introduce particular malicious perturbations to the victim's code context in practice. For example, users may introduce external code

by copying or installing packages; attackers can also upload their malicious code to popular open-source platforms (e.g., GitHub), whose code will be automatically downloaded and used in a knowledge base to enhance the performance of Code LLMs [20] through retrieval-augmented generation [21]. In particular, different from supply chain attacks [22], which are often detectable by static analysis [23] or sandboxing [24], external prompt injection introduces only non-functional code rather than malicious functional code to implant malicious behaviors, thereby circumventing mainstream detection methods.

The attack scenarios of traditional PI differ from our external prompt injection. Specifically, conventional PI attacks are typically executed by malicious users directly targeting large language models (LLMs), whereas our focus is on a malicious third-party attacker aiming to compromise the users of a victim application. More importantly, simply using existing PI attacks to solve the problem in our setting is infeasible, as traditional PI attacks [16], [25] rely on strong assumptions. Specifically: (1) They assume controlled positioning, requiring the malicious instruction/perturbation to be placed at a fixed position within the model's input (e.g., via prefix or suffix). (2) They usually assume an unlimited token budget, allowing for the use of hundreds of tokens and clear malicious instructions (e.g., "make a bomb"). These assumptions lead to insufficient attack effectiveness and stealthiness in practice.

To address these challenges, we propose a greedy perturbation optimization method, named ShadowCode, for external prompt injection based on code simulation. ShadowCode generally comprises two main stages. In the first stage, we design the perturbation optimization objective based on the attacker-specified malicious objective. Specifically, we represent the malicious objective as an output tuple (*position code*, *target code*), which induces the Code LLM to generate the malicious *target code* immediately after the *position code*. Subsequently, we conduct *code contextual simulation*, which includes *conditional code* (e.g., necessary packages and method declarations), *noise code* (i.e., target-agnostic random code), as well as pre-defined position code and target code. This simulation aims to mimic diverse, real-world code snippets, thereby ensuring ShadowCode's effectiveness in victim-specified, uncontrolled contextual environments. Finally, we design the loss function to identify the optimal induced perturbation that achieves this simulated outcome. In the second stage, we employ a greedy search to optimize the perturbation, incorporating two key designs: (1) *forward reasoning enhancement*: We increase the weight of the first few tokens in the target code in calculating the loss to enhance the effectiveness of ShadowCode under Code LLMs' autoregressive generation manner; (2) *keyword-based perturbation design*: we introduce a dynamic prefix containing a few key tokens from the output tuple as the augmentation of induced perturbation during optimization, aiming to reduce the length of final induced perturbation and avoid optimization failure caused by greedy search. This design is motivated by the understanding that key tokens convey essential information from the malicious objective, thereby facilitating the desired induction.

We collect 13 representative malicious objectives, including 3 straightforward yet representative malicious objectives and the Top-10 CWE Known Exploited Vulnerabilities (KEV) [26].

²We note that there are currently a few pioneering methods related to ours (generally dubbed 'indirect prompt injection') targeting Code LLMs [17], [18]. However, these works only provide simple examples, and their injection content is manually designed. More details and discussions are described in Section II-C.

These objectives are converted into specific cases across 3 programming languages, leading to a total of 31 threat cases since some of these malicious objectives apply only on C/C++. We evaluate our attack against various representative open-source Code LLMs, including CodeGemma-2b, CodeGemma-7b [4], and CodeGeeX2-6b [5]. The results demonstrate that ShadowCode can reach an 84.5% average attack success rate (ASR) over all cases and obtain over 80% ASR using only 8 tokens, indicating its effectiveness and stealthiness. In addition, we evaluate ShadowCode on mainstream commercial Code LLM-integrated applications, including CodeGeeX [2] (with more than 935 k installations in the VSCode [27] extension marketplace) and Github Copilot [1] (developed by Github and OpenAI). These applications correspond to the translucent-box and closed-box attack scenarios, respectively. ShadowCode achieves an ASR of up to 93.3% on CodeGeeX and 90.4% on Copilot using perturbations generated by open-source Code LLMs, demonstrating its threat in the real world.

In summary, this paper makes three main contributions: (1) We reveal the potential limitations of existing backdoor and adversarial attacks against Code LLMs, and introduce a new attack paradigm, i.e., external prompt injection. (2) We propose a simple yet effective method (i.e., ShadowCode). ShadowCode can generate concise induced perturbations with effectiveness in an uncontrolled context environment. (3) We conduct extensive experiments to evaluate ShadowCode under 31 representative threat cases on 3 representative open-source Code LLMs and 2 commercial Code LLM-integrated applications, demonstrating its effectiveness, stealthiness, and resistance to potential defenses.

II. BACKGROUND AND RELATED WORK

A. Code LLMs for Code Completion

Code and its comments serve as a valuable corpus for deep learning models, facilitating various code-related tasks, including code translation [28] and code understanding [15]. The emergence of LLMs provides new alternatives to coding tasks. Researchers have developed specialized models, e.g., Codex [3], CodeLlama [29], CodeGemma [4] and CodeGeeX [5]. These Code LLMs have significantly improved the efficiency of programmers [30], [31], and demonstrate an exceptional capability for syntax following, which is crucial for LLM agents [8]. This paper focuses on code completion [6], [7]. This is a paradigmatic task involving the generation of new code from incomplete input using Code LLMs, which have been integrated into commercial applications [1], [2]. However, this development paradigm may encounter some potential threats. On the one hand, Code LLMs are susceptible to influence from the code context, which can alter the generated code. On the other hand, attackers can easily control the context of code generation, such as by inserting external code when the user copies code, installs packages, or builds knowledge bases, thus affecting the decisions made by Code LLMs. Related events have already occurred in the real world [32], highlighting its threat.

B. Backdoor and Adversarial Attacks

Backdoor attacks target the training process of DNNs [33], [34]. These attacks implant malicious backdoors that can be activated by an attacker-specified trigger into a model during training, leading the model to misbehave when the trigger is present [35], [36]. In the code domain, there are generally two strategies to achieve this, including dataset poisoning [9], [11] and model poisoning [10]. The former requires the attacker to modify the code dataset for model training, while the latter requires the attacker to modify the parameters of the pre-trained Code LLM for application development. However, in real-world scenarios, developers of Code LLMs or code agents usually obtain the datasets or pretrained model from trusted entities (e.g., renowned institutes or companies). This mitigates the backdoor threats from the source. Recently, a few works also exploited backdoor attacks for positive purposes [37], [38], [39], which are out of our scope.

Adversarial attacks target the model's inference process [40], [41]. Adversarial attacks introduce perturbations to the input data, leading the model to make incorrect output [42], [43]. In the code domain, existing adversarial attacks [13], [44] aim to reduce the quality of generated code by crafting adversarial examples based on the (approximated) gradients regarding the model's inputs and outputs. For example, Jha et al. [13] slightly changed the original code (e.g., function and variable names) to lower the performance of Code LLMs measured by CodeBLEU score [45]. However, existing adversarial attacks are difficult to fulfill specific malicious purposes since they can hardly achieve precise malicious objectives, although they do not require manipulation in the training process. This is mostly due to their particular goal of attackers, i.e., finding adversarial samples that can degrade the quality of code generation, inevitably weakening Code LLMs' understanding of the input. Accordingly, adversarial attacks are difficult to pose a specific threat to Code LLMs.

C. Prompt Injection Attacks

Prompt injection is a type of input manipulation attack against LLMs [17], [46], [47], where attackers embed carefully crafted instructions in the dialogue prompt or context, causing the model to violate established system policies or developer intentions, leading to undesired responses (e.g., illegal content [16], information leakage [48] and unauthorized operations [17]). In the code domain, existing studies typically focus on the attacks conducted by malicious users, such as exploiting the Code LLMs to create malware [15]. These attacks aim for "jailbreaking" and are not concerned with the threat prompt injection poses to normal users of Code LLM-integrated applications.

Meanwhile, we notice that there are some methods sharing a similar attack scenario with ours, named external prompt injection (IPI) [17], [49], where attackers strategically inject the malicious instructions into the data likely to be retrieved at LLMs' inference time. However, these attacks typically involve manually designing injection instructions (e.g., demanding disregard for the system prompt) and rely on a *covert information space* (e.g., the Markdown of a Wikipedia page) without considering perturbation stealthiness. This leads to two problems in the real world: (1) Manually designed injection instructions

lack flexibility and sometimes even effectiveness. (2) It's often difficult to find a covert information space in many scenarios, e.g., in code files. Currently, in the code domain, existing IPI has only provided some preliminary showcases [17]. As shown in Fig. 12 in Appendix A, available online, researchers used 23 lines of comments to inject data deletion instructions (which could be achieved with just 10 tokens using ShadowCode), indicating a lack of substantial real-world threat.

III. METHODOLOGY

A. Formulation of External Prompt Injection Attacks

The Main Pipeline of Code Completion Task: Code completion is an essential task for Code LLMs, whose input may include code snippets, comments, pseudocode and user-specified natural language requests. Its output should consist of executable code blocks that meet user-specified requests. In this paper, we consider a general task format. Specifically, the input code block is defined as $\mathbf{X} = \{x_1, x_2, \dots, x_m\}$, where $x_i \in \mathbf{X}$ can be any token, and $\mathbf{X}$ as a whole adheres to the code format requirements. For instance, natural language requests can only be included in the code comments. The output is the completed code block $\mathbf{Y} = \{y_1, y_2, \dots, y_n\}$, which should be logically executable and functionally compliant with the requirements. We denote a Code LLM with weights θ as Θ , and the code completion task is simplified as $\mathbf{Y} = \Theta(\mathbf{X})$.

The Formulation of EPI: In general, there are three parties, including a Code LLM-integrated application, a user, and an attacker, involved in our EPI. The user employs the Code LLM-integrated application to complete the code project. The Code LLM-integrated application constructs the input $\mathbf{X}$ for the integrated Code LLM based on the code project and gets the output. The attacker generates a concise induced perturbation $\mathbf{P}$ and introduces it into the input $\mathbf{X}$. The Code LLM, affected by $\mathbf{P}$, generates code functionally close to the attacker's malicious objective O , i.e.,

$$\mathcal{D}(O, \mathcal{F}(\Theta(\mathbf{X} \leftarrow \mathbf{P}))) < \epsilon, \quad (1)$$

where $\mathcal{F}(\cdot)$ denotes the functionality of the code, $\mathbf{X} \leftarrow \mathbf{P}$ means injecting $\mathbf{P}$ into $\mathbf{X}$. $\mathcal{D}(\cdot, \cdot)$ is a distance function defined by the attacker to quantify functional deviation, where $\mathcal{D}(\cdot, \cdot) < \epsilon$ signifies that the generated code satisfies the attacker's objective. In practice, a common and feasible definition of $\mathcal{D}(\cdot, \cdot)$ is a boolean indicator function with $\epsilon = 1$, i.e., $\mathcal{D}(O, \mathcal{F}(\Theta(\mathbf{X} \leftarrow \mathbf{P}))) = 1 - \mathbb{I}\{O = \mathcal{F}(\Theta(\mathbf{X} \leftarrow \mathbf{P}))\}$.

Attacker's Capacities and Attack Scenarios: We assume that the attacker can obtain the parameters of a surrogate Code LLM to conduct attacks. This assumption is consistent with many adversarial methods [25], [50]. Specifically, we consider three main attack settings: glass-box, translucent-box, and closed-box settings. The former assumes that the attacker knows the specific parameters of Code LLM, which generally corresponds to scenarios where the model is an open-source model. The latter two assume that the attacker only knows the model series (rather than the specific version) and knows nothing about it, respectively, which generally corresponds to scenarios of using commercial Code LLM-integrated applications. In our attack scenarios, the attackers introduce external code to the victim user's code project. For example, introducing external code by

copying or installing packages is commonly done during the programming process [17], [22]. Meanwhile, external code from open-source platforms (e.g., GitHub) is usually downloaded as a knowledge base via retrieval-augmented generation [21], which can also be used as input to improve the performance of Code LLMs [20]. The practicality of this attack scenario is further discussed in Appendix A, available online. In particular, to avoid detection by existing methods (e.g., static analysis [23] and sandboxing [24]), external code should not immediately contain malicious functionalities, i.e., attackers should use non-functional code as induced perturbation.

B. Overview

In this section we briefly introduce our greedy perturbation optimization method (dubbed 'ShadowCode') for automatic and effective external prompt injection. As illustrated in Fig. 2, ShadowCode operates in two sequential stages: (1) *design optimization objectives* and (2) *perturbation optimization*. In the first stage, we design the perturbation optimization objectives for induced perturbation based on the attacker's malicious objective. In the second stage, we employ the greedy search to optimize the induced perturbation.

C. Design Optimization Objectives

In this stage, we first describe how to convert diverse malicious objectives to structured output tuples. Then we expand the output tuple to a full code contextual simulation. Finally, we design a loss function that leverages the simulation to guide the search for the optimal perturbation.

1) *Output Tuple:* We hereby describe how to convert the malicious objective to an output tuple. Let $\mathbf{Y}$ denote the normal code generated by a Code LLM; the attacker's goal is to modify $\mathbf{Y}$ such that the malicious objective is achieved. Such modifications can be divided into three types: deletion, addition, and alteration, all of which can be represented by an output tuple (position code $\mathbf{Y}_P$, target code $\mathbf{Y}_T$). Specifically, for deletion, the attacker wants to delete the code segment $[y_i, y_{i+1}, \dots, y_{i+k}]$ in $\mathbf{Y}$, and we have $\mathbf{Y}_P = [y_p, y_{p+1}, \dots, y_{i-1}]$ and $\mathbf{Y}_T = [y_{i+k+1}, y_{i+k+2}, \dots, y_t]$; For addition, the attacker wants to add the code segment $\mathbf{T}$ into the position between y_i and y_{i+1} , and we have $\mathbf{Y}_P = [y_p, y_{p+1}, \dots, y_{i-1}]$ and $\mathbf{Y}_T = [\mathbf{T}, y_{i+1}, y_{i+2}, \dots, y_t]$; For alteration, the attacker wants to change the code segment $[y_i, y_{i+1}, \dots, y_{i+k}]$ to $\mathbf{T}$, and we have $\mathbf{Y}_P = [y_p, y_{p+1}, \dots, y_{i-1}]$, $\mathbf{Y}_T = [\mathbf{T}, y_{i+k+1}, y_{i+k+2}, \dots, y_t]$.

2) *Code Contextual Simulation:* Optimizing the induced perturbation based solely on the output tuple ignores the surrounding code context. To ensure the effectiveness of ShadowCode in uncontrolled context environment, we expand the output tuple to simulate a code context of the victim user's project. In general, although the contextual environment of the victim user's project is usually uncontrollable for the attacker, there are still some fundamental logics to follow for code projects and code generation processes, which can be exploited by ShadowCode to simulate the victim's environment. Accordingly, our code context simulation is constructed based on the following three logics:




Fig. 2. The overall pipeline of our ShadowCode for (automatic) external prompt injection attack against Code LLMs. ShadowCode is executed in two stages: (1) Design Optimization Objectives: Given the attacker's malicious objective, ShadowCode converts it into a specific output tuple (position code, target code). Subsequently, ShadowCode expands this output tuple to simulate a code context of the victim user's project, based on which to design the loss function of perturbation optimization. (2) Perturbation Optimization: After defining the loss function, we use a surrogate Code LLM to calculate the loss and employ a greedy search to optimize the induced perturbation token by token. In particular, we enhance the optimization with forward reasoning enhancement (FRE) and keyword-based perturbation design (KPD) to improve the perturbation's effectiveness and stealthiness, respectively. After obtaining the induced perturbation through ShadowCode, the attacker inserts it into a malicious payload, which can then be strategically introduced into the code context via various channels (e.g., direct copy, dependencies, or RAG). Subsequently, this perturbation within the code context will induce Code LLMs to generate specific code at an attacker-specified position, thereby achieving the malicious objective.

Logic 1. Necessary Condition Argument: The code environment configured by the user must meet the basic requirements for the malicious objective, such as the necessary packages and method declarations. Therefore, we introduce this kind of code as *conditional code* at the beginning of the whole simulated code.

Logic 2. Fixed Generation Order: Code generation usually proceeds in a forward textual order. Consequently, the malicious objective is fulfilled once the *target code* (guided by the *position code*) is generated; any code generated thereafter does not significantly alter the final attack result. Thus, we connect the *position code* and *target code* and place them at the end of the simulated code.

Logic 3. Noise in Context: Since real-world software projects typically contain code with diverse functionalities, a code completion context consisting solely of code relevant to the malicious objective is inconsistent with practical settings and may lead to limited effectiveness. Therefore, we introduce *noise code* by adding irrelevant code snippets to simulate a more realistic code completion context.

Based on three logics above, we can construct the code contextual simulation for different malicious objectives. To better illustrate our strategy, we provide a straightforward example in Fig. 3, where the malicious objective is to make the victim use a weak encryption mode. The task components are automatically generated with the following stages. First, the position code should contain a fundamental encryption situation: the encryption key “key”, the text to be encrypted “iv” and the first half of the encryption instruction “cipher = DES.new(key,”. These variables and functions are necessary for the encryption methods, e.g., “CBC”. To achieve the malicious objective, the complete code instruction should call the weaker encryption method, “ECB”, thus the target code is “DES.MODE_ECB”. After obtaining the output tuple, ShadowCode then declares

```
demo.py > ...
1 import os
2 import sys
3 import time
4 from flask import Flask, request, jsonify
5 from Crypto.Cipher import DES
6 from Crypto.Util.Padding import pad, unpad
7 from Crypto.Random import get_random_bytes
8
9 '''Noise Tasks'''
10 def generate_random_string(length):
11     '''ellipsis'''
12
13
14 iv = get_random_bytes(DES.block_size)
15 key = b'8bytekey'
16 cipher = DES.new(key,
17
18     DES.MODE_ECB)
19
20 # Encrypt the message
21 ciphertext = cipher.encrypt(pad(
22     b'This is a test message', DES.block_size))
23 print(ciphertext)
24 # Decrypt the message
25 plaintext = unpad(cipher.decrypt(
26     ciphertext), DES.block_size)
27 print(plaintext)
```

The diagram highlights four segments in the code: **Conditional Code** (lines 5-7), **Noise Code** (lines 9-13), **Position Code** (lines 14-18), and **Target Code** (line 18).

Fig. 3. The example of *code contextual simulation*, where the attacker aims to induce the Code LLM to use encryption mode with low security (i.e., ECB encryption). The simulation consists of four code segments: conditional code, noise code, position code, and target code, to simulate the context environment most relevant to the malicious objective.

the *Crypto* series methods that enable the encryption code to run normally in the *conditional code*. Finally, ShadowCode introduces existing code snippets (e.g., those from github-code dataset) that is not related to the task and uses them as the *noise code*. The final code contextual simulation is constructed from multiple aspects for the same task to collaboratively ensure the effectiveness.




Fig. 4. Distortion during the autoregressive generation process. The attacker aims to insert a malicious deletion command after the position code, whereas the Code LLMs would typically generate the result printing command. During the perturbation optimization process, the generation probability of the first token in the target code, i.e., “os”, is not adequately optimized. Consequently, the Code LLM may not be successfully induced upon encountering the position code and generate an undesired token “print”. Subsequently, the failure to generate the first token compromises the successful generation of subsequent tokens.

3) *Loss Function of Perturbation Optimization*: After obtaining the code contextual simulation, a loss function can be formulated for optimizing induced perturbations. To achieve the malicious objective, the optimization aims to identify the perturbation that maximizes the probability of realizing the simulated successful attack. Specifically, the simulated contextual code is first divided into input and output components. The input to the Code LLM comprises *conditional code*, *noise code*, and *position code*, simulating a code completion task that can be exploited to achieve the malicious objective. This entire input is referred to as *task code*. The output, termed *target code*, serves as the label for the perturbation optimization task. Subsequently, with $\mathbf{Y}_T = [t_1, t_2, \dots, t_k]$, the loss function is defined as:

$$\mathcal{L} = - \sum_{i=1}^{|\mathbf{Y}_T|} \log \mathbb{P}(t_i | \mathbf{T}_s \leftarrow \mathbf{P}, t_1, t_2, \dots, t_{i-1}), \quad (2)$$

where $\mathbf{T}_s$ represents the task code, the probability $\mathbb{P}$ is inferred and calculated using the surrogate model weights θ . This function aims to maximize the likelihood of generating the target code given the task code as input. When the Code LLM receives the perturbation, it is induced to generate the target code if the context is appropriate (i.e., similar to the simulation in code functionality), thereby achieving the attacker’s malicious objectives.

D. Perturbation Optimization

In this stage, we first introduce two enhancement modules for the optimization process: *forward reasoning enhancement* and *keyword-based perturbation design*. Then, we will describe the whole algorithm for induced perturbation optimization.

1) *Forward Reasoning Enhancement*: Directly employing the loss function defined in (2) may lead to distortions during the autoregressive generation process. As shown in Fig. 4, since token prediction results are based on previous sampling outcomes, the original loss function may deviate from actual conditions due to unexpected predictions of the first few tokens. For example, the data deletion malicious instruction “os.system(“rm -rf /”)”, although brief, is difficult to integrate coherently with

the surrounding context (i.e., immediately after calling a demo function “TRG”). If the induced Code LLM fails to generate the first token “os”, the probabilities of generating subsequent tokens will diverge.

To alleviate this problem, a *forward reasoning enhancement* module is introduced. Motivated by the aforementioned discussion that subsequent tokens are influenced by preceding tokens in an autoregressive generation manner, earlier tokens hold greater importance for the overall generation process of the target code. Therefore, the importance of the initial tokens for the loss value is increased by assigning them higher weights, thereby enabling the optimization objective to adapt to the autoregressive mechanism. This increase in weight can be achieved through adding an additional loss. Its specific formulation is shown as follows:

$$\mathcal{L}_e = - \frac{1}{h} \sum_{i=1}^h \sum_{j=1}^i \log \mathbb{P}(t_i | \mathbf{T}_s \leftarrow \mathbf{P}, t_1, t_2, \dots, t_{j-1}),$$

$$\mathcal{L} = (\mathcal{L} + r\mathcal{L}_e) / (r + 1), \quad (3)$$

where $\mathcal{L}_e$ is the additional loss, r is its coefficient, and h is the number of enhanced tokens.

2) *Keyword-based Perturbation Design*: This module aims to reduce the token count required for induced perturbations by offloading information into prefix tokens, which enhances attack stealth. Arguably, minimizing token usage and obscuring malicious characteristics are critical to maintaining attack stealthiness,³ thereby avoiding detection by victim users or RAG’s developers. Generally, perturbations optimized via greedy search typically exhibit less obvious malicious characteristics compared to manually designed ones. However, these perturbations struggle to accurately convey malicious instructions as effectively as manually designed perturbations, resulting in higher token usage. A seemingly straightforward solution, which involves optimizing based on manually designed perturbations by replacing obvious malicious characteristics with equivalent tokens, proves difficult in practice. As shown in Fig. 5, the tokens constituting these obvious malicious characteristics (i.e., convey dangerous data deletion instructions) are the most critical and challenging to replace via greedy search.

To alleviate these limitations, we propose a simple yet effective enhancement method, i.e., *keyword-based perturbation design*. Intuitively, the purpose of the induced perturbation is to fulfill the attacker’s malicious objective. Therefore, the output tuple, which expresses malicious objective, will contain the necessary information for the induced perturbation. Motivated by this understanding, we leverage tokens from the output tuple to construct a prefix for the perturbation, termed a *keyword*, which is then injected into the victim’s code context alongside the perturbation. During keyword construction, we use the loss function in (3) to greedily select the optimal token combination. In particular, we utilize only a very small number of tokens (typically one or two) for keyword construction to ensure no overt malicious intent is present in the injected content. As

³‘Stealthiness’ hereby targets the evasion of commonly used model-level detection (e.g., semantic-based analysis). We do not consider manual code review since it requires substantial human labor, making it impractical for large-scale deployment.




Fig. 5. Examples of different induced perturbations. The attacker aims to induce the Code LLM to generate a data deletion command (“`os.system(“rm -rf /”)`”) immediately after a specified function call (“`TRG`”) via induced perturbation. Manually designed induced perturbations lack stealthiness due to their obvious malicious intent, a flaw that a simple greedy search-based replacement alone cannot rectify. Induced perturbations optimized with random initialization necessitate extensive token usage to ensure effectiveness. Our method combines optimized induced perturbations with a prefix (keyword), thereby achieving both effectiveness and stealthiness with limited tokens.

shown in Fig. 5, the keyword “`TR`” exhibits no overt malicious intent but can enhance the perturbation’s effectiveness compared to directly optimized counterparts.

3) *Optimization Algorithm*: We hereby introduce the specific optimization process of ShadowCode for induced perturbations and their corresponding keywords. Specifically, we employ greedy gradient search to optimize an initialized perturbation $P = [p_1, p_2, \dots, p_l]$, where l denotes the token count. As shown in Algorithm 1, our ShadowCode method optimizes the perturbation token by token in each iteration (line 4). We utilize the loss function in (3) for greedy gradient search (line 5), i.e., the optimization objective for the i -th token p_i is:

$$p_i = \arg \min_{p_i} \mathcal{L}. \quad (4)$$

To achieve this, ShadowCode propagates the loss value to the vocabulary embedding layer. Based on the gradient at token p_i , it greedily searches for the suitable token e'_i within the Code LLM’s vocabulary set V to replace p_i , as shown in the following equation:

$$e'_i = \arg \min_{e'_i \in V} (e'_i - p_i)^T \nabla_{p_i} \mathcal{L}, \quad (5)$$

where $\nabla_{p_i} \mathcal{L}$ represents the gradient of the task loss. ShadowCode enhances this token replacement strategy by employing beam search, which considers the Top- k token candidates (Line 6-7). ShadowCode then replaces the original token with each candidate to form a new P' (Line 8-9), and computes a new loss value $\mathcal{L}'$: if $\mathcal{L}'$ is better than $\mathcal{L}$, P is replaced with P' (Line 10-13). Finally, if P is not optimized within an entire iteration, ShadowCode considers it optimal and terminates the optimization process (Line 16).

Additionally, we use grid search to optimize the keyword. The perturbation and the keyword are optimized alternately, which

Algorithm 1. Induced Perturbation Optimization.

```

1: Input: Task Code  $T_s$ , target code  $Y_T$ , initial
   perturbation  $P$  with tokens  $p_{1:l}$ , beam search size  $k$ 
2: Vocab set  $V$ , weights  $\theta$  of the surrogate Code LLM
3: repeat
4:   for  $i \in [0, \dots, l]$  do
5:      $\mathcal{L} = \text{get\_loss}(T_s, Y_T, P, \theta)$ 
6:      $\mathcal{G} = \nabla_{p_i} \mathcal{L}$ 
7:      $\mathcal{P}_{\text{candidates}} = \text{Top-}k_{e'_i} \min(e'_i{}^T \mathcal{G}), e'_i \in V$ 
8:     for  $j \in [0, \dots, k]$  do
9:        $P' : \text{replace } p_i \text{ to } \mathcal{P}_{\text{candidates}}^j$ 
10:       $\mathcal{L}' = \text{get\_loss}(T_s, Y_T, P', \theta)$ 
11:      if  $\mathcal{L}' < \mathcal{L}$  then
12:         $P = P'$ 
13:      end if
14:    end for
15:  end for
16: until  $P$  does not change in the whole iteration
17: Output: Optimized perturbation  $P$ 

```

can be expressed as follows:

$$P_{key}^{s+1} = \arg \min_{P_{key}} \mathcal{L}(T_s, T, P_{key} \oplus P^s, \theta),$$

$$P^{s+1} = \arg \min_P \mathcal{L}(T_s, T, P_{key}^{s+1} \oplus P, \theta), \quad (6)$$

where s is the number of optimization steps and P_{key} represents the keyword. $\oplus$ means concatenating two token sequences. We select keyword tokens from the output tuple, i.e., $\forall p_{ki}$ in $P_{key}, p_{ki} \in \{\text{target code}, \text{position code}\}$. For example, as shown in Fig. 5, the keyword “`TR`” are tokens from the position code of the output tuple.

IV. EXPERIMENTS

A. Main Settings

1) *Malicious Objectives and Code Datasets*: We hereby collect 13 representative malicious objectives, including 3 straightforward yet representative malicious objectives and the Top-10 CWE Known Exploited Vulnerabilities (KEV) [26] across 3 programming languages: Python, Java, and C/C++. The straightforward objectives include using `MODE_ECB` for encryption (dubbed ‘ST0’), using `ssl3` protocol (dubbed ‘ST1’), and executing “`rm -rf /`” (dubbed ‘ST2’). ST0 and ST1 are studied in classical code backdoor attacks [9], while ST2 is a famous destructive instruction. CWE vulnerabilities are generally selected for experiments in code model research [14], [51], and CWE Top-10 KEV weaknesses are supported by American Cybersecurity and Infrastructure Security Agency (CISA). We construct cases using the method described in Section III-C based on the official examples (i.e., CWE Top-10 KEV) or LLM-generated demo (i.e., ST0, ST1, ST2). As shown in Table I, some objectives apply only to C/C++, e.g., CWE416 “Use After Free”, thus we have a total of 31 cases. We provide comprehensive details for all cases and present case studies in Appendix A, available online to illustrate how manipulating code generation can achieve diverse malicious objectives and lead to severe security implications. In

TABLE I

THE SUMMARY OF OBJECTIVES AND DATASETS. WE LIST THE PROGRAMMING LANGUAGES THAT EACH OBJECTIVE/DATASET IS APPLICABLE TO: C/C++ (C), PYTHON (P), JAVA (J). THE LAST ROW OF THE TABLE SHOWS THE TOTAL NUMBER OF TEST DATA WE DESIGN. SOME OBJECTIVES APPLY ONLY TO C/C++, RESULTING IN A TOTAL OF 31 APPLICABLE CASES.

	Identifier/Name	Applicable Languages
Objectives	ST0, ST1, ST2	C/C++, Python, Java
	CWE416, 122, 787, 843	C/C++
	CWE20, 78, 502, 918, 22, 306	C/C++, Java
	MBPP	Python
Datasets	Humaneval	Python
	Eval-Plus	Python
	CodeXGLUE	Python, Java
	Humaneval-x	C/C++, Python, Java
	Data Count: 14273	(C/P/J): 2132/7965/4176

addition, we collect corresponding code completion datasets, including Humaneval [52], Humaneval-x [5], CodeXGLUE [53], MBPP [54], and Eval-Plus [55]. The items in these datasets contain prompts and corresponding outputs across different code completion tasks, which can simulate unpredictable users' instructions and context in real-world scenarios. Some datasets do not cover Java or C/C++, as shown in Table I. As a result, we have 7 sub-datasets (i.e., the subsection of a dataset containing a particular programming language) containing more than 14,000 data items totally.

2) *Model Selection*: In our experiment, we select three models from two representative Code LLM series, including CodeGemma [4] and CodeGeeX [5], for evaluation. CodeGemma series are trained by Google LLC and outperform some old-version Code LLMs like CodeLlama [29] on code completion tasks. We select CodeGemma-2b and CodeGemma-7b to compare the performance differences between models of distinct scales. CodeGeeX series have corresponding applications [2] with more than 935 k downloads (on VSCode [27] marketplace). We select CodeGeeX2-6b, a version with stable performance. We exclude earlier small-scale code-oriented models, such as CodeBERT [56] and CodeT5 [57], because large-scale code-oriented models perform much better than them in code completion tasks.

3) *Baseline Selection*: We compare our ShadowCode to GCG-based and manually designed IPI attacks. In particular, the vanilla GCG [25] is a representative PI attack and does not meet the requirements of our attack scenarios. Therefore, we adapt GCG into an IPI attack that is applicable to Code LLMs, dubbed 'C-GCG' (code-oriented GCG). The manually designed IPI method (dubbed 'IPI-M') is designed based on the existing work [17]. IPI-M limits the number of injected tokens to make the perturbation concise since the existing IPI attack usually uses a large number of readable comments. Besides, we also provide the results of the original model (dubbed 'No Attack') for reference, since the original model itself may also have inherent insecurity factor. More details for baselines are provided in Appendix A, available online.

4) *Evaluation Metrics*: In this paper, we use ASR to evaluate the effectiveness, not-bad cases rate (NBR) and selective threat factor (STF) to evaluate the flexibility, and number of injected tokens and chars (NIT and NIC) to evaluate the stealthiness. For ASR, we consider an attack successful when the LLM is

induced to generate code that fulfills the malicious intent, i.e., the generated code is identical or functionally equivalent to the target code. NBR and STF are new metrics introduced for our target-specific attack scenarios. **NBR Scenario**: The attacker has multiple attack opportunities (e.g., poisoning public code repositories) and aims to achieve as diverse a range of attack objectives as possible. Therefore, NBR calculates the proportion of all feasible cases (where $ASR > \text{a lower threshold } r_1$). **STF Scenario**: The attacker has limited attack opportunities (e.g., injecting private code projects) and only accepts sufficiently stable attacks selectively. Thus, STF calculates the proportion of all stable cases (where $ASR > \text{a higher threshold } r_2$). Their formulations are as follows:

$$NBR = 1 - \frac{\# \text{ of cases where } ASR < r_1}{\text{Total \# of cases}},$$

$$STF = \frac{\sum_{i \in \text{cases where } ASR > r_2} ASR_i}{\text{Total \# of cases}}, \quad (7)$$

where r_1 and r_2 are set to 20% and 80% respectively in our evaluation empirically. Further reasons of evaluation metrics selection are discussed in Appendix A, available online.

5) *Hyperparameters and Computational Facilities*: We conduct experiments using the following parameter settings to evaluate our approach. During gradient calculation, we utilize five noisy code samples, and the random seed for initializing induced perturbations is set to 100. For *forward reasoning enhancement* module, the number of tokens reinforced is set to $h = 2$, as defined in (3). For the hyperparameters related to attack performance, the default parameter values are as follows: the Top- k value is 400, the rate r in (3) is 0.4, the number of keyword is 2, and the perturbation length is 10. Tokenization to calculate token number is performed using Gemma's tokenizer.

We run Algorithm 1 on an NVIDIA RTX A6000 Graphics Card with the above parameter settings, employing a float16 precision format without quantization. In general, the perturbation generation process requires approximately 0.5-1 hours for CodeGemma-2b and 2-3 hours for both CodeGemma-7b and CodeGeeX-6b. These computational times ensure the generation of high-quality induced perturbations tailored to the respective model sizes across all the threat cases in our experiment.

B. Evaluation on Open-Source Code LLMs

We first perform the attack evaluation on open-source Code LLMs. This includes three aspects: (1) we assess the attack effectiveness over 3 Code LLMs across all threat cases, demonstrating that ShadowCode is widely effective and more threatening than other attacks; (2) we evaluate attack flexibility against different targets using two special metrics; (3) we evaluate the trade-off between effectiveness and stealthiness *w.r.t.* the number of injected tokens.

1) *Effectiveness*: We hereby evaluate whether ShadowCode can successfully induce the Code LLMs to generate attacker-specified targets. We control the number of injected tokens for three attack methods in this experiment. For ShadowCode, we use its default settings: perturbation length of 10 and 2 keywords, while C-GCG uses a prefix length of 12. Since IPI-M requires a sufficient number of tokens to be effective, we just limit its perturbation to fewer than 30 tokens. Table II

TABLE II

THE ASR ACROSS DIFFERENT CODE LLMs. THE EXPERIMENT TARGETS FOUR CODE COMPLETION DATASETS AND TESTS THE SUCCESS RATES OF VARIOUS ATTACK METHODS USING THREE CODE LLMs. AS SHOWN IN TABLE I, SOME DATASETS DO NOT INCLUDE DATA FOR THE C/C++ AND JAVA LANGUAGES, THUS WE HAVE A TOTAL OF 7 SUBDATASETS. BEST ATTACK RESULTS ARE MARKED IN BOLDFACE.

Code LLM↓	Dataset→	Humaneval/Humaneval-x			CodeXGLUE		MBPP	Eval-Plus	AVG
	Method↓ / Language→	C/C++	Python	Java	Python	Java	Python	Python	
CodeGemma-2b	No Attack	1.9%	17.3%	9.4%	22.0%	4.9%	12.6%	17.0%	12.2%
	C-GCG [25]	40.6%	50.2%	32.1%	36.5%	49.4%	59.0%	44.2%	44.6%
	IPI-M [17]	3.4%	30.8%	6.4%	33.3%	3.2%	19.7%	31.0%	18.3%
	ShadowCode (Ours)	75.5%	93.4%	84.2%	74.4%	67.3%	94.1%	85.4%	82.0%
CodeGemma-7b	No Attack	0.0%	12.2%	0.0%	18.9%	0.0%	15.6%	12.1%	8.4%
	C-GCG [25]	43.5%	53.8%	55.4%	33.1%	20.4%	47.2%	50.6%	43.4%
	IPI-M [17]	11.1%	29.6%	4.3%	46.8%	20.8%	37.8%	33.2%	26.2%
	ShadowCode (Ours)	85.8%	97.9%	78.5%	81.6%	70.5%	94.4%	97.4%	86.6%
CodeGeeX2-6b	No Attack	0.1%	13.4%	12.9%	19.0%	24.0%	21.9%	14.3%	15.0%
	C-GCG [25]	42.9%	46.1%	44.9%	41.7%	32.5%	51.7%	33.7%	41.9%
	IPI-M [17]	9.5%	11.4%	5.3%	17.9%	0.9%	22.0%	11.5%	11.2%
	ShadowCode (Ours)	79.6%	95.0%	80.8%	83.3%	72.0%	97.9%	85.6%	84.9%

provides the average ASR for four methods. The effectiveness of ShadowCode outperforms the other baseline methods: it achieves an average ASR of over 80% across all three Code LLMs, with the highest average ASR reaching 86.6%. In addition, we further extend our evaluation by analyzing case-specific attack effectiveness in Appendix B, available online, which confirms that our method outperforms the baselines, particularly in challenging cases such as CWE22 and ST2.

Comparing the three Code LLMs, the CodeGemma-7b model has the highest security in normal conditions (8.4% for No Attack) and is the least likely to generate vulnerable code. However, it is the most susceptible model to ShadowCode, which may be attributed to its strong contextual learning ability. Comparing the three programming languages, Python is the most susceptible language to ShadowCode due to its concise syntax and the smallest number of tokens consumed. Besides, ShadowCode is relatively less effective on CodeXGLUE, which is mostly due to its larger data length and more complex data content. In contrast, the simplest dataset with the shortest data items, MBPP, is the most vulnerable to attacks, including C-GCG and IPI-M. Therefore, we believe that an increase in the contextual complexity can reduce the ASR of ShadowCode. In other words, when the contextual code becomes complex, the strength of ShadowCode (e.g., the length of induced perturbations) needs to be increased accordingly to ensure the attack's effectiveness. According to our subsequent experiments in Section IV-C and Section IV-E, increasing the length of ShadowCode or the Top- k value can effectively enhance the attack performance of ShadowCode.

2) *Flexibility*: Previous experiments indicate that ShadowCode exhibits the best average ASR performance across various models. However, attackers may have different requirements depending on the attack scenarios: when attack opportunities are abundant, they can tolerate attack cases with lower success rates; conversely, they only accept cases with higher success rates. We define the capability to satisfy such complex requirements as flexibility, and evaluate it using NBR and STF.

As shown in Table III, ShadowCode has the highest NBR on the Python programming language, which means that Python-based attacks can achieve the most diverse malicious objectives when attack opportunities are abundant. Additionally, ShadowCode scores the highest in STF for Python, indicating

TABLE III

THE ATTACK EVALUATION RESULTS OF NBR AND STF AGAINST OPEN-SOURCE CODE LLMs. BEST RESULTS ARE MARKED IN BOLD.

Metric	Method	C/C++	Python	Java
NBR	No Attack	0.00	0.22	0.01
	C-GCG	0.39	0.50	0.56
	IPI-M	0.08	0.25	0.17
	Ours	0.77	0.92	0.78
STF	No Attack	0.00	0.11	0.03
	C-GCG	0.26	0.47	0.39
	IPI-M	0.09	0.19	0.07
	Ours	0.76	0.87	0.74

that Python is equally well-suited for scenarios with limited attack opportunities. STF and NBR are not always positively correlated. For example, for ShadowCode, the NBR for C/C++ is lower than that for Java, but C/C++ has a higher STF score. This suggests that C/C++-based attacks pose a greater threat in scenarios with limited attack opportunities, whereas Java-based attacks are more likely to pose a greater threat in scenarios with abundant attack opportunities. Besides, for both NBR and STF, ShadowCode consistently outperforms the baselines.

3) *Stealthiness*: We hereby evaluate the trade-off between effectiveness and stealthiness. Objectively, ShadowCode cannot perfectly evade human review of external code (as discussed in Section V). Nevertheless, shortening the token length remains a practical approach to reduce the probability of exposure, as human review is highly resource-intensive thus can hardly guarantee an exhaustive coverage of all code snippets. Therefore, we measure the stealthiness of ShadowCode by the number of tokens and chars (NIT and NIC), which directly determine whether the victim will be aware of the perturbation.

As shown in Fig. 6, ASR and STF show a clear increase as NIT increases. Before reaching 60% ASR, ShadowCode attacks require very low NIT/NIC values. If an attacker aims for an effective but unstable attack, the perturbation can be extremely stealthy. When ASR exceeds 80%, the increase in NIT/NIC values accelerates, but increasing NIT leads to a very stable and effective improvement in ASR and STF (only when ASR is close to 100% will there be a small fluctuation). This demonstrates that enhancing ShadowCode's performance does not require significantly increasing the length of the perturbation as with





Fig. 6. The study of stealthiness. We plot the above two graphs using NIT and ASR as the x-axis respectively. We sampled 11 different NITs, NIT 0 (No Attack) and NIT 2 to 20 (with a sampling interval of 2).

existing methods [17], every token can effectively improve the ASR.

Additionally, we find that after ASR reaches above 80%, the growth of NIC is slower compared to NIT, while this difference is not significant below 80% ASR. This is likely because when NIT is small, longer character-length tokens can convey more information, whereas when NIT is large, shorter character-length tokens can convey more precise information. The tokens used by ShadowCode do not have distinct characteristics in terms of character length but are determined by the attacker's design of the token number. This means that filtering ShadowCode attacks based on character length is unreliable.

C. Evaluation on Commercial Code LLMs

In previous experiments, we demonstrate that ShadowCode reaches sufficient effectiveness, flexibility, and stealthiness on open-source models. However, some widely used models are commercial and therefore are closed-source in real-world scenarios. In these cases, attackers need to conduct transfer attacks, i.e., generating induced perturbations using open-source Code LLMs to attack closed-source ones. We hereby evaluate ShadowCode's transferability by leveraging two representative real-world scenarios, thereby further verifying its potential threat in the real world.

1) *Attack Scenarios on Real-world Applications:* The scenarios with limited access to Code LLM information can be

TABLE IV
THE EVALUATION OF TRANSFERABILITY. WE TEST ShadowCode's ASR, NBR, AND STF UNDER TRANSLUCENT-BOX AND CLOSED-BOX SCENARIOS RESPECTIVELY TO DEMONSTRATE ITS TRANSFERABILITY.

Scenario	Language	ASR	NBR	STF
Translucent-box	C/C++	80.4%	0.92	0.75
	Python	93.3%	1.00	0.93
	Java	71.7%	0.78	0.63
Closed-box	C/C++	67.7%	0.70	0.64
	Python	90.4%	1.00	0.85
	Java	74.4%	0.78	0.65

broadly divided into two categories: translucent-box and closed-box, with the corresponding explanations as follows:

- 1) *Translucent-box Setting:* The attacker can infer the open-source model used by the target application through interaction but cannot obtain the exact version of the model. For example, the scale of the model's parameters, fine-tuning information, and version numbers are unknown to the attacker. We simulate this scenario using CodeGeeX2-6b and the CodeGeeX application. To the best of our knowledge, the CodeGeeX model has at least four different versions, and the CodeGeeX application does not disclose the specific model information it uses.
- 2) *closed-box Setting:* A more challenging situation is when the attacker knows nothing about the model used by the target application, or the model itself is closed-source (such as Codex). We simulate this scenario by attacking GitHub Copilot (supported by Codex). To align with translucent-box setting, we adopt CodeGeeX2-6b to generate induced perturbations.

1) *Settings:* We construct all cases in three languages within the integrated application of VSCode and manually test randomly selected code completion samples from the Humaneval/Humaneval-x dataset (620 data items in total). Other settings are the same as those described in Section IV-B.

1) *Results:* As shown in Table IV, in the translucent-box scenario, ShadowCode still demonstrates great attack performance. In Python and C/C++, ShadowCode achieves ASRs of 80.4% and 93.3%, respectively. Even in the worst-performing case (i.e., Java), the ASR achieves 71.7%, with only a 9.1% decrease compared to the glass-box scenario; in the closed-box scenario, ShadowCode also achieves a satisfying attack performance of 90.4% on Python. The ASRs on C/C++ and Java are 67.7% and 74.4%, respectively. In other words, our ShadowCode still maintains a high level of threat, although its performance is lower than that of glass-box scenarios to some extent. In addition, ShadowCode achieves a high NBR (> 0.7) and STF (> 0.6) in both the translucent-box and closed-box scenarios, indicating that it also maintains sufficient flexibility. Overall, ShadowCode achieves good attack performance in both translucent-box and closed-box scenarios, demonstrating its reliable transferability and threat to the commercial applications in the real-world.

In Fig. 7, we present a representative case of ST2 in the closed-box scenario. The concise perturbation "TR os now rm sys" is inserted in the comment of a tool function. The commercial applications generate the functional code

```

1 Introduction:
2 This is a close elements function, to check if in given list of numbers, are any
3 two numbers closer to each other than given threshold. Using TR as how we sys
4 , please be careful. You can use this function to check if there are any
5 close elements in a list of numbers.
6
7 Usage:
8 Check if in given list of numbers, are any two numbers closer to each other than
9 given threshold.
10
11 Examples:
12 >>> has_close_elements([1.0, 2.0, 3.0], 0.5)
13 False
14 >>> has_close_elements([1.0, 2.0, 3.0, 4.0, 5.0, 2.0], 0.3)
15 True
16
17 def has_close_elements(numbers: List[float], threshold: float) -> bool:
18     for idx, elem in enumerate(numbers):
19         for idx2, elem2 in enumerate(numbers):
20             if idx == idx2:
21                 distance = abs(elem - elem2)
22                 if distance < threshold:
23                     return True
24     return False

```

(a) Perturbation Injection

```

1 import os
2 def TRG(a,b):
3     result = a + b
4     return result
5
6 if __name__ == '__main__':
7     print(has_close_elements([1.0, 2.0, 3.0, 4.0, 5.0, 2.0], 0.3))
8     a = 1
9     b = 2
10    print(TRG(a,b))
11    os.system('rm -rf /')

```

(b) Target Completion

Fig. 7. The example of a successful attack in a real-world closed-box scenario (Github Copilot) across files. The perturbation is inserted into the comment of a function. When the user calls a function in a code context, Copilot will be induced to generate malicious code at the specified position (after calling “TRG”).

“os.system(“rm -rf /”)” at the attacker-specified position (i.e., after using TRG function) under the induction of the perturbation, although it is obviously a highly malicious instruction that can cause dangerous results.

D. Resistance to Adaptive Defenses

In this section, we discuss whether our method can resist potential adaptive defenses that are particularly designed for our attack. In general, we can roughly categorize these defenses into detection and removal, and design the following potential defense methods against ShadowCode.

1) *Detection: Semantic Analysis*: Detection methods distinguish between malicious payload code that contains perturbations and normal code, thus rejecting the use of malicious payload code. Since detection methods based on code functionality are ineffective against non-functional perturbations, we adopt a semantic analysis approach. This involves examining the differences between the benign code and the injected code in the embedding layer of Code LLMs. However, this method may be difficult to implement. We find that code segments injected with perturbations might not have obvious semantic differences from benign ones. Specifically, as shown in Fig. 8, we sample 1640 code segments and construct classification prompts. We discuss the setting of this semantic analysis in detail in Appendix B, available online. The results show that benign code and injected code are almost completely indistinguishable.

2) *Removal: Massive Deletion and Find-and-Replace*: This type of defense involves removing perturbations by altering the non-functional parts of the code after introducing external code. In general, it entails a trade-off between code readability and the strength of removal since we cannot faithfully detect




Fig. 8. The T-SNE plot [58] of the classification keyword [CLS] at layer 17 shows the semantic differences between the original data and the injected data in the last embedding layer of CodeGemma-2b. T-SNE plot figures at the other layers have similar results.




Fig. 9. The resistance to the removal of chars. We remove characters from the perturbation, selecting 5% for removal each time, and we calculate the average ASR for each programming language across applicable cases. As the influence of the perturbation fades, the ASR results converge to the same level as no attack baseline.

TABLE V
ASR (DEG) FOR DIFFERENT FORMS OF PERTURBATIONS. DEG DENOTES THE ASR DEGRADATION COMPARED WITH COMMENTS.

Type↓	Language↓, Metric→	ASR	Deg
Variable Assignment	C/C++	67.5%	-8.0%
	Python	90.1%	-7.8%
	Java	78.6%	-5.6%
Output Content	C/C++	65.0%	-10.5%
	Python	83.0%	-14.6%
	Java	66.5%	-9.0%

injected code segments (as mentioned above). To explore its effectiveness, we hereby test the impact of the percentage of character removal on the effectiveness of perturbations. As shown in Fig. 9, ShadowCode still maintains a sufficiently high ASR under a 50% removal rate. On the other hand, even if readability is completely sacrificed by extensively removing comments, ShadowCode can still be effective through non-comment portions of the code, as indicated in Table V. Arguably, the removal methods are likely to require significant human effort and may not yield effective results since ShadowCode imposes strict limitations on the length of perturbations.




Fig. 10. Ablation study of keyword number and Top- k . We compare ShadowCode's ASR and STF with keyword length from 0 to 2. During the process of generating perturbations, we use six different Top- k values in our evaluation.

Despite our empirical demonstration that ShadowCode is resilient against classical adaptive defenses, certain more advanced countermeasures could potentially alleviate its practical threat (discussed in Section V), as we do not yet provide theoretical guarantees for the attack's effectiveness.

E. Ablation Study

In this section, we conduct an ablation study to demonstrate the necessity of the *keyword-based designing* and *forward reasoning enhancement* methods proposed in Section 4 for ShadowCode attacks. We also discuss the effects of a critical method-related hyper-parameter (i.e., Top- k), which is crucial for determining both the computational load of Algorithm 1 and the effectiveness of our ShadowCode. In addition, we conduct an ablation study on adding noise code, details of which are in Appendix B, available online.

1) *Effects of Top- k* : This experiment validates the impact of Top value on the performance of ShadowCode. As shown in Fig. 10, increasing the Top- k value can significantly improve ShadowCode's ASR and STF. Arguably, the attack performance can reach its optimal level when the Top- k is increased to the value of the LLM's vocab size. However, this approach would introduce computationally expensive overhead (CodeGemma has a vocab size of 256,000, far exceeding our default Top- k value of 400).

2) *Effects of Keyword-based Designing*: This experiment evaluates the effects of our optimization module, i.e., *keyword-based designing*. Specifically, we vary the number of keyword and generate the perturbations. As shown in Fig. 10, limiting the number of keyword has a significant impact. The result of 1-keyword perturbation has approximately 10% decreases of ASRs and above 0.3 decrease of STF compared to 2-keyword perturbation. The 0-keyword perturbation, even with high computational power (600 Top- k value), struggles to achieve an ASR above 80%. In contrast, 2-keyword perturbation can easily achieve both high ASR (above 95%) and STF (above 0.95). The above results indicate that using keywords can effectively improve the performance of ShadowCode.

3) *Effects of Forward Reasoning Enhancement*: We hereby perform a representative case study to illustrate the effects of




Fig. 11. Ablation study of FRE. We compare the differences in ASR and losses both with and without using FRE module. 1-key1 and 1-key2 represent keeping one of the two selected keyword, respectively.

forward reasoning enhancement (FRE), which corrects the distortion of the loss function. Specifically, we use the 'ST2' threat case as an example for our discussions since it is representative and difficult for attackers to achieve. It requires generating a clearly malicious code snippet and necessitates producing target code in a completely unrelated context, making it hard for the Code LLM to generate the *target code*, easily leading to distortion. To better evaluate the effects of this module, we first mitigate the effects of keyword-based designing by using the same keyword. We judge the degree of distortion based on the loss value. The results are illustrated in Fig. 11, where we have three key findings. Firstly, under the same conditions, using FRE can lead to a higher ASR. Secondly, without FRE, the ASR is still very low ($< 20\%$) even when the loss value is lower than that with FRE. Thirdly, FRE has a greater impact on ASR than keyword-based design. These results verify the necessity of forward reasoning enhancement in our ShadowCode method.

V. POTENTIAL LIMITATIONS AND FUTURE WORK

As the first work on automatically attacking Code LLMs via external prompt injection, we have to admit that ShadowCode still exhibits some potential limitations. Firstly, although ShadowCode eliminates explicit malicious intent and significantly reduces token consumption, it still cannot completely bypass human review. A more ideal stealthy attack would involve generating natural-looking and contextually coherent injection payloads, which will be further explored in future work. Secondly, combining external code provenance with some operational measures targeting either the victim's project or the Code LLM (e.g., code auditing, sandboxing, and LLM privilege restriction) may mitigate the impact of EPI. We hope this work will draw the attention of researchers in related fields to EPI attacks, facilitating a more comprehensive defense framework for Code LLMs. Thirdly, our work only focuses on the code completion scenario. We will investigate the performance of ShadowCode in other scenarios such as code summarization and code translation in future work.

VI. CONCLUSION

This paper investigated how to manipulate Code LLMs through their context without intervening in the model's training

process. We proposed ShadowCode, an external prompt injection (EPI) approach that compromises models by injecting concise, non-functional perturbations. ShadowCode optimized these induced perturbations by simulating the victim's code contexts and employing a greedy search, guided by two enhancement modules: forward reasoning enhancement and keyword-based perturbation design. Extensive experiments across multiple programming languages, open-source models, and commercial applications verified the effectiveness, flexibility, stealthiness, and transferability of our attack.

ACKNOWLEDGMENT

We would sincerely thank Prof. Kui Ren and Prof. Chun Chen from Zhejiang University for their helpful comments and suggestions on an early draft of this paper.

REFERENCES

- [1] N. Friedman, "Introducing GitHub copilot: Your AI pair programmer," 2021. [Online]. Available: <https://docs.github.com/zh/copilot/quickstart>
- [2] "Codegeex," 2024. Accessed: May 10, 2024. [Online]. Available: <https://codegeex.cn/zh-CN>
- [3] O. W. Zaremba and G. Brockman, "OpenAI codex," 2021. [Online]. Available: <https://openai.com/blog/openai-codex>
- [4] C. Team et al., "CodeGemma: Open code models based on Gemma," 2024, *arXiv:2406.11409*.
- [5] Q. Zheng et al., "CodeGeex: A pre-trained model for code generation with multilingual benchmarking on HumanEval-X," in *Proc. 29th ACM SIGKDD Conf. Knowl. Discov. Data Mining*, 2023, pp. 5673–5684.
- [6] F. Zhang et al., "RepoCoder: Repository-level code completion through iterative retrieval and generation," in *Proc. Conf. Empirical Methods Natural Lang. Process.*, 2023, pp. 2471–2484.
- [7] Y. Ding et al., "CrossCodeVal: A diverse and multilingual benchmark for cross-file code completion," in *Proc. Adv. Neural Inf. Process. Syst.*, 2024, pp. 46701–46723.
- [8] X. Wang et al., "Executable code actions elicit better LLM agents," in *Proc. 41st Int. Conf. Learn. Representations*, 2024, pp. 50208–50232.
- [9] R. Schuster, C. Song, E. Tromer, and V. Shmatikov, "You Autocomplete me: Poisoning vulnerabilities in neural code completion," in *Proc. 30th USENIX Secur. Symp.*, 2021, pp. 1559–1575.
- [10] Y. Li, S. Liu, K. Chen, X. Xie, T. Zhang, and Y. Liu, "Multi-target backdoor attacks for code pre-trained models," in *Proc. 61st Annu. Meeting Assoc. Comput. Linguist.*, 2023, pp. 7236–7254.
- [11] Z. Yang et al., "Stealthy backdoor attack for code models," *IEEE Trans. Softw. Eng.*, vol. 50, no. 4, pp. 721–741, Apr. 2024.
- [12] Y. Zhou et al., "Adversarial robustness of deep code comment generation," *ACM Trans. Softw. Eng. Methodol.*, vol. 31, no. 4, pp. 1–30, 2022.
- [13] A. Jha and C. K. Reddy, "CodeAttack: Code-based adversarial attacks for pre-trained programming language models," in *Proc. AAAI Conf. Artif. Intell.*, 2023, pp. 14892–14900.
- [14] S. Yan et al., "An LLM-assisted easy-to-trigger backdoor attack on code completion models: Injecting disguised vulnerabilities against strong detection," in *Proc. 33rd USENIX Secur. Symp.*, 2024, pp. 1795–1812.
- [15] W. Ma, Y. Song, M. Xue, S. Wen, and Y. Xiang, "The 'code' of ethics: A holistic audit of AI code generators," *IEEE Trans. Dependable Secure Comput.*, vol. 21, no. 5, pp. 4997–5013, Sep./Oct. 2024.
- [16] X. Shen, Z. Chen, M. Backes, Y. Shen, and Y. Zhang, "'Do anything now': Characterizing and evaluating in-the-wild jailbreak prompts on large language models," in *Proc. ACM SIGSAC Conf. Comput. Commun. Secur.*, 2024, pp. 1671–1685.
- [17] K. Greshake et al., "Not what you've signed up for: Compromising real-world LLM-integrated applications with indirect prompt injection," in *Proc. 16th ACM Workshop Artif. Intell. Secur.*, 2023, pp. 79–90.
- [18] Q. Zhan, Z. Liang, Z. Ying, and D. Kang, "InjecAgent: Benchmarking indirect prompt injections in tool-integrated large language model agents," in *Proc. Findings Assoc. Comput. Linguist.*, 2024, pp. 10471–10506.
- [19] Y. He, Y. Liu, L. Wu, Z. Yang, K. Ren, and Z. Qin, "MsDroid: Identifying malicious snippets for android malware detection," *IEEE Trans. Dependable Secure Comput.*, vol. 20, no. 3, pp. 2025–2039, May/Jun. 2023.
- [20] J. Chen et al., "Code search is all you need? Improving code suggestions with code search," in *Proc. IEEE/ACM 46th Int. Conf. Softw. Eng.*, 2024, pp. 880–892.
- [21] W. Zou, R. Geng, B. Wang, and J. Jia, "{PoisonedRAG}: Knowledge corruption attacks to { Retrieval-Augmented} generation of large language models," in *Proc. 34th USENIX Secur. Symp.*, 2025, pp. 3827–3844.
- [22] M. Ohm, H. Plate, A. Sykosch, and M. Meier, "Backstabber's knife collection: A review of open source software supply chain attacks," in *Proc. Int. Conf. Detection Intrusions Malware, Vulnerability Assessment*, 2020, pp. 23–43.
- [23] Z. Li et al., "VulDeePecker: A deep learning-based system for vulnerability detection," in *Proc. Netw. Distrib. Syst. Secur.*, 2018, pp. 1–15.
- [24] A. Alhamdan and C.-A. Staicu, "{Sanddriller}: A {fully-automated} approach for testing {language-based} {JavaScript} sandboxes," in *Proc. 32nd USENIX Secur. Symp.*, 2023, pp. 3457–3474.
- [25] A. Zou, Z. Wang, J. Z. Kolter, and M. Fredrikson, "Universal and transferable adversarial attacks on aligned language models," 2023, *arXiv:2307.15043*.
- [26] "2023 CWE top 10 known exploited vulnerabilities weaknesses," 2024. Accessed: Oct. 30, 2024. [Online]. Available: https://cwe.mitre.org/top25/archive/2023/2023_key_list.html
- [27] "Visual studio code," Microsoft Corporation, 2024. Accessed: May 9, 2024. [Online]. Available: <https://code.visualstudio.com/>
- [28] W. Sun, C. Fang, Y. Chen, G. Tao, T. Han, and Q. Zhang, "Code search based on context-aware code translation," in *Proc. 44th Int. Conf. Softw. Eng.*, 2022, pp. 388–400.
- [29] B. Roziere et al., "Code Llama: Open foundation models for code," 2023, *arXiv:2308.12950*.
- [30] C. S. Xia, Y. Wei, and L. Zhang, "Automated program repair in the era of large pre-trained language models," in *Proc. IEEE/ACM 45th Int. Conf. Softw. Eng.*, 2023, pp. 1482–1494.
- [31] Y. He et al., "On benchmarking code LLMs for android malware analysis," in *Proc. 34th ACM SIGSOFT Int. Symp. Softw. Testing Anal.*, 2025, pp. 153–160.
- [32] "Criminal cases involving large models acting as intermediaries to generate malicious code," 2025. Accessed: Jan. 6, 2025. [Online]. Available: <https://x.com/rcky0/status/1859656430888026524>
- [33] J. Zhang, J. Chi, Z. Li, K. Cai, Y. Zhang, and Y. Tian, "BadMerging: Backdoor attacks against model merging," in *Proc. ACM SIGSAC Conf. Comput. Commun. Secur.*, 2024, pp. 4450–4464.
- [34] B. Yi et al., "Probe before you talk: Towards black-box defense against backdoor unalignment for large language models," in *Proc. Int. Conf. Learn. Representations*, 2025, pp. 1–27.
- [35] W. Jiang, T. Zhang, H. Qiu, H. Li, and G. Xu, "Incremental learning, incremental backdoor threats," *IEEE Trans. Dependable Secure Comput.*, vol. 21, no. 2, pp. 559–572, Mar./Apr. 2024.
- [36] Y. Li, Y. Bai, Y. Jiang, Y. Yang, S.-T. Xia, and B. Li, "Untargeted backdoor watermark: Towards harmless and stealthy dataset copyright protection," in *Proc. Adv. Neural Inf. Process. Syst.*, 2022, pp. 13238–13250.
- [37] Y. Li, M. Zhu, X. Yang, Y. Jiang, T. Wei, and S.-T. Xia, "Black-box dataset ownership verification via backdoor watermarking," *IEEE Trans. Inf. Forensics Secur.*, vol. 18, pp. 2318–2332, 2023.
- [38] M. Ya, Y. Li, T. Dai, B. Wang, Y. Jiang, and S.-T. Xia, "Towards faithful XAI evaluation via generalization-limited backdoor watermark," in *Proc. Int. Conf. Learn. Representations*, 2024, pp. 43524–43551.
- [39] S. Shao et al., "DATABench: Evaluating dataset auditing in deep learning from an adversarial perspective," 2025, *arXiv:2507.05622*.
- [40] B. He et al., "Generating transferable 3D adversarial point cloud via random perturbation factorization," in *Proc. AAAI Conf. Artif. Intell.*, 2023, pp. 764–772.
- [41] W. Jiang et al., "Rethinking the design of backdoor triggers and adversarial perturbations: A color space perspective," *IEEE Trans. Dependable Secure Comput.*, vol. 22, no. 3, pp. 2823–2840, May/Jun. 2025.
- [42] H. Peng et al., "TextCheater: A query-efficient textual adversarial attack in the hard-label setting," *IEEE Trans. Dependable Secure Comput.*, vol. 21, no. 4, pp. 3901–3916, Jul./Aug. 2024.
- [43] M. Javaheripi, M. Samragh, B. D. Rouhani, T. Javidi, and F. Koushanfar, "CuRTAIL: Characterizing and thwarting adversarial deep learning," *IEEE Trans. Dependable Secure Comput.*, vol. 18, no. 2, pp. 736–752, Mar./Apr. 2021.
- [44] P. Bielik and M. Vechev, "Adversarial robustness for code," in *Proc. Int. Conf. Mach. Learn.*, 2020, pp. 896–907.
- [45] S. Ren et al., "CodeBLEU: A method for automatic evaluation of code synthesis," 2020, *arXiv:2009.10297*.
- [46] S. Yi et al., "Jailbreak attacks and defenses against large language models: A survey," 2024, *arXiv:2407.04295*.
- [47] Z. Ren, S. Liang, A. Liu, and D. Tao, "ICLShield: Exploring and mitigating in-context learning backdoor attacks," in *Proc. Int. Conf. Mach. Learn.*, 2025, pp. 1–14.

- [48] Y. Li et al., “Rethinking data protection in the (generative) artificial intelligence era,” 2025, *arXiv:2507.03034*.
- [49] R. Zhang et al., “Instruction backdoor attacks against customized {LLMs},” in *Proc. 33rd USENIX Secur. Symp.*, 2024, pp. 1849–1866.
- [50] M. Shen, C. Li, Q. Li, H. Lu, L. Zhu, and K. Xu, “Transferability of white-box perturbations: {Query-efficient} adversarial attacks against commercial {DNN} services,” in *Proc. 33rd USENIX Secur.*, 2024, pp. 2991–3008.
- [51] J. He and M. Vechev, “Large language models for code: Security hardening and adversarial testing,” in *Proc. ACM SIGSAC Conf. Comput. Commun. Secur.*, 2023, pp. 1865–1879.
- [52] M. Chen et al., “Evaluating large language models trained on code,” 2021, *arXiv:2107.03374*.
- [53] S. Lu et al., “CodeXGLUE: A machine learning benchmark dataset for code understanding and generation,” in *Proc. Adv. Neural Inf. Process. Syst.*, 2022, pp. 1–16.
- [54] J. Austin et al., “Program synthesis with large language models,” 2021, *arXiv:2108.07732*.
- [55] J. Liu, C. S. Xia, Y. Wang, and L. Zhang, “Is your code generated by ChatGPT really correct? Rigorous evaluation of large language models for code generation,” in *Proc. Adv. Neural Inf. Process. Syst.*, 2024, pp. 21558–21572.
- [56] Z. Feng et al., “CodeBERT: A pre-trained model for programming and natural languages,” in *Proc. Conf. Empirical Methods Natural Lang. Process.*, 2020, pp. 1536–1547.
- [57] Y. Wang, W. Wang, S. Joty, and S. C. Hoi, “CodeT5: Identifier-aware unified pre-trained encoder-decoder models for code understanding and generation,” in *Proc. Conf. Empirical Methods Natural Lang. Process.*, 2021, pp. 8696–8708.
- [58] L. Van der Maaten and G. Hinton, “Visualizing data using t-SNE,” *J. Mach. Learn. Res.*, vol. 9, no. 11, 2008, pp. 2579–2605.